

Accelerating K-Means Clustering with OpenMP and GPU Parallelization

Final project

Master's program in Informatics Engineering and Internet of Things (IoT)

Student: Darya Martsinouskaya – 26610

Course Unit: High Performance Computing

Teacher: José Jasnaú Caeiro

<div>01</div> <div>Introduction</div> <p>K-means is a widely used clustering algorithm in machine learning.</p> <p>It partitions datasets into k clusters by minimizing intra-cluster distances.</p> <p>The computational cost grows significantly for large datasets [1].</p> <p>The work explores parallelization strategies using OpenMP and GPU acceleration to improve efficiency.</p>	<div>02</div> <div>Problem description</div> <p>Sequential K-means: The standard algorithm iteratively assigns points to clusters and updates centroids.</p> <p>Challenges: High time complexity limits scalability for large datasets (e.g., MNIST).</p> <p>Goal: Reduce execution time while maintaining accuracy.</p>	<div>03</div> <div>Parallelization Strategies</div> <p>OpenMP (Multicore CPU): Parallelization of main loops for point assignment and centroid updates.</p> <p>GPU Acceleration: Offloading computations to GPU for enhanced parallel execution.</p> <p>Speedup Calculation: Speedup (s) = Time (Sequential) / Time (Parallel)</p>	<div>04</div> <div>Experimental Setup</div> <p>Dataset: MNIST (grayscale images, 28x28 pixels).</p> <p>Metrics: Execution time, accuracy, speedup.</p> <p>Implementations:</p> <ul style="list-style-type: none">Sequential (baseline)OpenMP parallelizedGPU-accelerated																													
<div>05</div> <div>Experimental Setup</div> <p>Dataset: MNIST (grayscale images, 28x28 pixels).</p> <p>Metrics: Execution time, accuracy, speedup.</p> <p>Implementations:</p> <ul style="list-style-type: none">Sequential (baseline)OpenMP parallelizedGPU-accelerated	<div>06</div> <div>Results</div> <table><tr><th>k</th><th>Sequential Time (s)</th><th>OpenMP Time (s)</th><th>GPU Time (s)</th><th>Speedup OpenMP</th><th>Speedup GPU</th></tr><tr><td>10</td><td>838.213</td><td>100.817</td><td>1.26322</td><td>8.31</td><td>663.55</td></tr><tr><td>50</td><td>1495.33</td><td>48.6627</td><td>1.23969</td><td>30.73</td><td>1206.21</td></tr><tr><td>100</td><td>921.122</td><td>20.0596</td><td>1.2539</td><td>45.92</td><td>734.61</td></tr><tr><td>200</td><td>691.691</td><td>40.2045</td><td>1.24277</td><td>17.20</td><td>556.57</td></tr></table> <ul style="list-style-type: none">GPU acceleration significantly outperforms both sequential and OpenMP implementations, achieving up to 1206.21× speedupOpenMP provides moderate improvements (up to 45.92× speedup) but loses efficiency at higher k valuesAll implementations maintain identical clustering accuracy, confirming correctness	k	Sequential Time (s)	OpenMP Time (s)	GPU Time (s)	Speedup OpenMP	Speedup GPU	10	838.213	100.817	1.26322	8.31	663.55	50	1495.33	48.6627	1.23969	30.73	1206.21	100	921.122	20.0596	1.2539	45.92	734.61	200	691.691	40.2045	1.24277	17.20	556.57	<div>07</div> <div>Conclusion</div> <ul style="list-style-type: none">Implemented sequential, OpenMP parallel, and GPU-accelerated versions of K-meansOpenMP parallelization achieved up to 45.92× speedup, but gains decreased for large cluster sizesGPU acceleration provided the highest improvement, reaching 1206.21× speedup with excellent scalabilityAccuracy remained consistent across all implementations, validating parallelization correctnessFurther optimizations (e.g., CUDA, memory management) could enhance performance for larger datasets
k	Sequential Time (s)	OpenMP Time (s)	GPU Time (s)	Speedup OpenMP	Speedup GPU																											
10	838.213	100.817	1.26322	8.31	663.55																											
50	1495.33	48.6627	1.23969	30.73	1206.21																											
100	921.122	20.0596	1.2539	45.92	734.61																											
200	691.691	40.2045	1.24277	17.20	556.57																											

References

- AnjaTaniovićandVukVranjković.“ImplementationofparallelK-meansalgorithmfor image classification using OpenMP and MPI libraries”. In: 2024 Zooming Innovation in Consumer Technologies Conference (ZINC). 2024, pp. 54–59. doi: 10.1109/ZINC61849.2024.10579351.